

shopify_dsci_2021_fall

January 20, 2022

```
[1]: import pandas as pd
```

```
[2]: df = pd.read_csv("shopify_sneakers.csv")
```

0.1 Getting Familiar with the Data

```
[3]: df.head()
```

```
[3]:   order_id  shop_id  user_id  order_amount  total_items  payment_method \
0          1       53      746        224             2           cash
1          2       92      925         90             1           cash
2          3       44     861        144             1           cash
3          4       18     935        156             1  credit_card
4          5       18     883        156             1  credit_card

              created_at
0  2017-03-13 12:36:56
1  2017-03-03 17:38:52
2  2017-03-14 4:23:56
3  2017-03-26 12:43:37
4  2017-03-01 4:35:11
```

```
[4]: df.dtypes
```

```
[4]: order_id        int64
shop_id         int64
user_id         int64
order_amount    int64
total_items     int64
payment_method  object
created_at      object
dtype: object
```

This is some minor data cleaning and to just confirm that the data set is in a 30-day period

```
[5]: df["created_at"] = pd.to_datetime(df['created_at'])
dateSorted = df.sort_values(by = ["created_at"])
```

```
dateSorted.head()
```

```
[5]:      order_id  shop_id  user_id  order_amount  total_items  payment_method \
1862          1863       39      738         536            4        cash
1741          1742       39      910         268            2        cash
3228          3229       97      912         324            2        cash
1267          1268       80      798         290            2  credit_card
2689          2690       49      799         258            2  credit_card

           created_at
1862 2017-03-01 00:08:09
1741 2017-03-01 00:10:19
3228 2017-03-01 00:14:12
1267 2017-03-01 00:19:31
2689 2017-03-01 00:22:25
```

```
[6]: print("First\n", dateSorted.iloc[0])
print("Last\n", dateSorted.iloc[len(df) - 1])
```

```
First
order_id              1863
shop_id                39
user_id                738
order_amount            536
total_items               4
payment_method          cash
created_at        2017-03-01 00:08:09
Name: 1862, dtype: object
Last
order_id              2458
shop_id                95
user_id                700
order_amount            168
total_items                 1
payment_method        credit_card
created_at        2017-03-30 23:55:35
Name: 2457, dtype: object
```

0.2 Question 1a

Assuming the AOV is a simple calculation of: $\frac{\text{total_amount}}{\text{total_orders}}$

Then this would probably just be a mean of the `order_amount` column.

```
[7]: df["order_amount"].mean()
```

```
[7]: 3145.128
```

As outlined in the challenge, this is much higher than expected if we consider the fact that we are dealing with sneaker shops. If we take a look at the dataset itself to see what is going on:

```
[8]: df["order_amount"].max()
```

```
[8]: 704000
```

The AOV is so high because there is a purchase order worth \$704,000!

```
[9]: df.iloc[df["order_amount"].argmax()]
```

```
[9]: order_id          16
shop_id            42
user_id           607
order_amount      704000
total_items       2000
payment_method    credit_card
created_at        2017-03-07 04:00:00
Name: 15, dtype: object
```

And now we understand why, this looks to be a bulk order of some kind, and assuming there are several of these in the data set, this is why the AOV is skewed higher than expected.


0.3 Questions 1b and 1c

The problem with using an average or a mean calculation in data analysis, is that they are easily influeunced by outliers. In the context of our problem, an outlier could be the example above: a bulk order for 2000 sneakers. Rather than using a mean, I would recommend taking a trimmed mean instead, which can reduce the influence from the extreme outliers we have in our data but still hopefully calculate a reasonable AOV.

The question however is what percentage to trim by.

```
[10]: import matplotlib.pyplot as plt
plt.scatter(df.order_id, df.order_amount)
```

```
[10]: <matplotlib.collections.PathCollection at 0x7f50d3456d60>
```



I count roughly 30 or so values that seem much too high to be normal from a visual glance which is around 0.6% however just to be a bit more assured, we'll trim by 1% instead.

```
[11]: from scipy import stats

amounts = df.order_amount
percentage = 1/100
stats.trim_mean(amounts, percentage).round(2)
```

[11]: 372.16

With a slight adjustment to the AOV calculation, we now have a more reasonable value of \$372.16

0.3.1 Question 2a

```
[ ]: """
SELECT COUNT(OrderID) FROM [Orders]
WHERE ShipperID = (SELECT ShipperID FROM [Shippers]
WHERE ShipperName = "Speedy Express")
```

ANSWER: 54

0.3.2 Question 2b

```
[ ]: """
SELECT LastName FROM [Employees]
WHERE EmployeeID = (SELECT EmployeeID FROM [Orders]
GROUP BY EmployeeID
ORDER BY COUNT(*) DESC
LIMIT 1)
```

ANSWER: Peacock

```
"""
```

0.3.3 Question 2c

```
[ ]: """
SELECT ProductName FROM [Products]
WHERE ProductID = (SELECT ProductID FROM [OrderDetails]
WHERE OrderID IN (SELECT OrderID FROM [Orders]
WHERE CustomerID IN (SELECT_
↳CustomerID FROM [Customers]
↳WHERE Country = "Germany"))
GROUP BY ProductID
ORDER BY SUM(Quantity) DESC
LIMIT 1)
```

ANSWER: Boston Crab Meat

```
"""
```

□